

MATH-564 Project

House Prices Data Analytics

A20495939 - Jasleen Kaur Bhatia

A20504279 - Sajesh Rao Erabelli

Problem Statement

House price varies based on the condition of itself and the environment. From the number of bedrooms to the location of the house, any variable might be the key that affects the house price the most. In this project , we will use ANOVA and MLR to determine the relation of house situations with sold price and predict the house price.

Loading Libraries

```
library(readr)
library(stringr)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(caTools)
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(ggplot2)
library(grid)
library(lattice)
library(gridExtra)
```

```
##
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dplyr':  
##  
##      combine
```

```
library(ggpubr)  
library(tidyverse)
```

```
## — Attaching packages  
## —————  
## tidyverse 1.3.2 —
```

```
## ✓ tibble  3.1.8      ✓ purrr   0.3.5  
## ✓ tidyr   1.2.1      ✓ forcats 0.5.2  
## — Conflicts ————— tidyverse_conflicts() —  
## ✖ gridExtra::combine() masks dplyr::combine()  
## ✖ dplyr::filter()      masks stats::filter()  
## ✖ dplyr::lag()         masks stats::lag()
```

```
library(broom)  
library(AICcmodavg)  
library(caret)
```

```
##  
## Attaching package: 'caret'  
##  
## The following object is masked from 'package:purrr':  
##  
##      lift
```

```
library(leaps)  
library(MASS)
```

```
##  
## Attaching package: 'MASS'  
##  
## The following object is masked from 'package:dplyr':  
##  
##      select
```

```
library(car)
```

```
## Loading required package: carData
##
## Attaching package: 'car'
##
## The following object is masked from 'package:purrr':
##
##     some
##
## The following object is masked from 'package:dplyr':
##
##     recode
```

Loading Datasets

```
house_datasales <- read_csv("/Users/jasleenkaurbhatia/Desktop/Semester3/Applied_Stats/AS
Project/kc_house_data.csv")
```

```
## Rows: 21597 Columns: 21
## — Column specification —————
## Delimiter: ","
## chr (1): date
## dbl (20): id, price, bedrooms, bathrooms, sqft_living, sqft_lot, floors, wat...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
head(house_datasales)
```

```
## # A tibble: 6 × 21
##       id date      price bedro...1 bathr...2 sqft_...3 sqft_...4 floors water...5 view
##       <dbl> <chr>      <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl> <dbl>
## 1 7129300520 10/13/... 2.22e5     3     1     1180    5650     1     0     0
## 2 6414100192 12/9/2... 5.38e5     3    2.25    2570    7242     2     0     0
## 3 5631500400 2/25/2... 1.8 e5     2     1     770    10000     1     0     0
## 4 2487200875 12/9/2... 6.04e5     4     3     1960    5000     1     0     0
## 5 1954400510 2/18/2... 5.1 e5     3     2     1680    8080     1     0     0
## 6 7237550310 5/12/2... 1.23e6     4    4.5    5420   101930     1     0     0
## # ... with 11 more variables: condition <dbl>, grade <dbl>, sqft_above <dbl>,
## #   sqft_basement <dbl>, yr_built <dbl>, yr_renovated <dbl>, zipcode <dbl>,
## #   lat <dbl>, long <dbl>, sqft_living15 <dbl>, sqft_lot15 <dbl>, and
## #   abbreviated variable names 1bedrooms, 2bathrooms, 3sqft_living, 4sqft_lot,
## #   5waterfront
```

```
colnames(house_datasales)
```

```
## [1] "id"           "date"         "price"        "bedrooms"
## [5] "bathrooms"    "sqft_living"  "sqft_lot"     "floors"
## [9] "waterfront"   "view"         "condition"    "grade"
## [13] "sqft_above"   "sqft_basement" "yr_built"     "yr_renovated"
## [17] "zipcode"      "lat"          "long"         "sqft_living15"
## [21] "sqft_lot15"
```

```
dim(house_datsales)
```

```
## [1] 21597    21
```

```
str(house_datsales)
```

```
## spc_tbl_ [21,597 × 21] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ id      : num [1:21597] 7.13e+09 6.41e+09 5.63e+09 2.49e+09 1.95e+09 ...
## $ date    : chr [1:21597] "10/13/2014" "12/9/2014" "2/25/2015" "12/9/2014" ...
## $ price   : num [1:21597] 221900 538000 180000 604000 510000 ...
## $ bedrooms : num [1:21597] 3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms : num [1:21597] 1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_living : num [1:21597] 1180 2570 770 1960 1680 ...
## $ sqft_lot   : num [1:21597] 5650 7242 10000 5000 8080 ...
## $ floors    : num [1:21597] 1 2 1 1 1 1 2 1 1 2 ...
## $ waterfront : num [1:21597] 0 0 0 0 0 0 0 0 0 0 ...
## $ view      : num [1:21597] 0 0 0 0 0 0 0 0 0 0 ...
## $ condition : num [1:21597] 3 3 3 5 3 3 3 3 3 3 ...
## $ grade     : num [1:21597] 7 7 6 7 8 11 7 7 7 7 ...
## $ sqft_above : num [1:21597] 1180 2170 770 1050 1680 ...
## $ sqft_basement : num [1:21597] 0 400 0 910 0 1530 0 0 730 0 ...
## $ yr_built  : num [1:21597] 1955 1951 1933 1965 1987 ...
## $ yr_renovated : num [1:21597] 0 1991 0 0 0 ...
## $ zipcode   : num [1:21597] 98178 98125 98028 98136 98074 ...
## $ lat       : num [1:21597] 47.5 47.7 47.7 47.5 47.6 ...
## $ long      : num [1:21597] -122 -122 -122 -122 -122 ...
## $ sqft_living15 : num [1:21597] 1340 1690 2720 1360 1800 ...
## $ sqft_lot15  : num [1:21597] 5650 7639 8062 5000 7503 ...
## - attr(*, "spec")=
## .. cols(
## ..   id = col_double(),
## ..   date = col_character(),
## ..   price = col_double(),
## ..   bedrooms = col_double(),
## ..   bathrooms = col_double(),
## ..   sqft_living = col_double(),
## ..   sqft_lot = col_double(),
## ..   floors = col_double(),
## ..   waterfront = col_double(),
## ..   view = col_double(),
## ..   condition = col_double(),
## ..   grade = col_double(),
## ..   sqft_above = col_double(),
## ..   sqft_basement = col_double(),
## ..   yr_built = col_double(),
## ..   yr_renovated = col_double(),
## ..   zipcode = col_double(),
## ..   lat = col_double(),
## ..   long = col_double(),
## ..   sqft_living15 = col_double(),
## ..   sqft_lot15 = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

So the total number of rows in housedata dataset is : 21597 and number of columns is 21.

Understanding the data :

```
summary(house_datasales)
```

```
##           id           date           price           bedrooms
## Min.      :1.000e+06   Length:21597   Min.       : 78000   Min.       : 1.000
## 1st Qu.:2.123e+09   Class :character   1st Qu.: 322000   1st Qu.: 3.000
## Median :3.905e+09   Mode  :character   Median : 450000   Median : 3.000
## Mean      :4.580e+09           Mean      : 540297   Mean      : 3.373
## 3rd Qu.:7.309e+09           3rd Qu.: 645000   3rd Qu.: 4.000
## Max.      :9.900e+09           Max.      :7700000   Max.      :33.000
##   bathrooms   sqft_living   sqft_lot   floors
## Min.      :0.500   Min.      : 370   Min.      : 520   Min.      :1.000
## 1st Qu.:1.750   1st Qu.: 1430   1st Qu.: 5040   1st Qu.:1.000
## Median :2.250   Median : 1910   Median : 7618   Median :1.500
## Mean      :2.116   Mean      : 2080   Mean      : 15099   Mean      :1.494
## 3rd Qu.:2.500   3rd Qu.: 2550   3rd Qu.: 10685   3rd Qu.:2.000
## Max.      :8.000   Max.      :13540   Max.      :1651359   Max.      :3.500
##   waterfront   view   condition   grade
## Min.      :0.000000   Min.      :0.0000   Min.      :1.00   Min.      : 3.000
## 1st Qu.:0.000000   1st Qu.:0.0000   1st Qu.:3.00   1st Qu.: 7.000
## Median :0.000000   Median :0.0000   Median :3.00   Median : 7.000
## Mean      :0.007547   Mean      :0.2343   Mean      :3.41   Mean      : 7.658
## 3rd Qu.:0.000000   3rd Qu.:0.0000   3rd Qu.:4.00   3rd Qu.: 8.000
## Max.      :1.000000   Max.      :4.0000   Max.      :5.00   Max.      :13.000
##   sqft_above   sqft_basement   yr_built   yr_renovated
## Min.      : 370   Min.      : 0.0   Min.      :1900   Min.      : 0.00
## 1st Qu.:1190   1st Qu.: 0.0   1st Qu.:1951   1st Qu.: 0.00
## Median :1560   Median : 0.0   Median :1975   Median : 0.00
## Mean      :1789   Mean      : 291.7   Mean      :1971   Mean      : 84.46
## 3rd Qu.:2210   3rd Qu.: 560.0   3rd Qu.:1997   3rd Qu.: 0.00
## Max.      :9410   Max.      :4820.0   Max.      :2015   Max.      :2015.00
##   zipcode   lat   long   sqft_living15
## Min.      :98001   Min.      :47.16   Min.      : -122.5   Min.      : 399
## 1st Qu.:98033   1st Qu.:47.47   1st Qu.: -122.3   1st Qu.:1490
## Median :98065   Median :47.57   Median : -122.2   Median :1840
## Mean      :98078   Mean      :47.56   Mean      : -122.2   Mean      :1987
## 3rd Qu.:98118   3rd Qu.:47.68   3rd Qu.: -122.1   3rd Qu.:2360
## Max.      :98199   Max.      :47.78   Max.      : -121.3   Max.      :6210
##   sqft_lot15
## Min.      : 651
## 1st Qu.: 5100
## Median : 7620
## Mean      : 12758
## 3rd Qu.: 10083
## Max.      :871200
```

```
for (column in house_datasales){
  print( typeof(column))
}
```

```
## [1] "double"
## [1] "character"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
## [1] "double"
```

So we understand that we have 21 features and all but one have their datatype as double. Only one specific feature- date has the data type as "character".

It would be better if we create a dataset without the values of date as that will allow us to understand the data better by using correlation and other functions/plots.

```
house_datsales1 <- house_datsales[,-1:-2]
colnames(house_datsales1)
```

```
## [1] "price"      "bedrooms"   "bathrooms"  "sqft_living"
## [5] "sqft_lot"   "floors"     "waterfront" "view"
## [9] "condition"  "grade"      "sqft_above"  "sqft_basement"
## [13] "yr_built"   "yr_renovated" "zipcode"     "lat"
## [17] "long"       "sqft_living15" "sqft_lot15"
```

```
dim(house_datsales1)
```

```
## [1] 21597    19
```

```
#View(house_datsales1)
```

Loading USZIPCODE data:

```
zipcode_data <- read_csv("/Users/jasleenkaurbhatia/Desktop/Semester3/Applied_Stats/AS Project/uszip.csv")
```

```
## Rows: 33121 Columns: 18
## — Column specification —————
## Delimiter: ","
## chr (10): zip, city, state_id, state_name, county_fips, county_name, county...
## dbl (4): lat, lng, population, density
## lgl (4): zcta, parent_zcta, imprecise, military
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
head(zipcode_data)
```

```
## # A tibble: 6 × 18
##   zip      lat    lng city  state...1 state...2 zcta  paren...3 popul...4 density count...5
##   <chr> <dbl> <dbl> <chr> <chr>    <chr>    <lgl> <lgl>      <dbl>    <dbl> <chr>
## 1 00601  18.2 -66.8 Adjun... PR      Puerto... TRUE   NA        17113    103.  72001
## 2 00602  18.4 -67.2 Aguada PR      Puerto... TRUE   NA        37751    476  72003
## 3 00603  18.5 -67.1 Aguad... PR      Puerto... TRUE   NA        47081    575.  72005
## 4 00606  18.2 -66.9 Maric... PR      Puerto... TRUE   NA         6392    58.3  72093
## 5 00610  18.3 -67.1 Anasco PR      Puerto... TRUE   NA        26686    287.  72011
## 6 00612  18.4 -66.7 Areci... PR      Puerto... TRUE   NA        59369    339.  72013
## # ... with 7 more variables: county_name <chr>, county_weights <chr>,
## #   county_names_all <chr>, county_fips_all <chr>, imprecise <lgl>,
## #   military <lgl>, timezone <chr>, and abbreviated variable names 1state_id,
## #   2state_name, 3parent_zcta, 4population, 5county_fips
```

```
colnames(zipcode_data)
```

```
## [1] "zip"           "lat"           "lng"           "city"
## [5] "state_id"      "state_name"    "zcta"          "parent_zcta"
## [9] "population"    "density"       "county_fips"   "county_name"
## [13] "county_weights" "county_names_all" "county_fips_all" "imprecise"
## [17] "military"      "timezone"
```

```
dim(zipcode_data)
```

```
## [1] 33121    18
```

So, uszips dataset have 33788 rows and 18 columns.

Understanding the USZIPS data :

```
summary(zipcode_data)
```



```

##      zip                lat                lng                city
## Length:33121      Min.      :-14.22      Min.      :-176.63      Length:33121
## Class :character      1st Qu.: 35.39      1st Qu.: -97.23      Class :character
## Mode  :character      Median : 39.49      Median : -88.19      Mode  :character
##                               Mean  : 38.82      Mean   : -90.92
##                               3rd Qu.: 42.12      3rd Qu.: -80.22
##                               Max.   : 71.27      Max.    : 145.75
##
##      state_id          state_name          zcta          parent_zcta
## Length:33121      Length:33121      Mode:logical      Mode:logical
## Class :character      Class :character      TRUE:33121      NA's:33121
## Mode  :character      Mode  :character
##
##
##
##      population          density          county_fips          county_name
## Min.      :      0      Min.      :      0.0      Length:33121      Length:33121
## 1st Qu.:      707      1st Qu.:      7.6      Class :character      Class :character
## Median :      2804      Median :      30.5      Mode  :character      Mode  :character
## Mean   :      9910      Mean   :      509.7
## 3rd Qu.:     13481      3rd Qu.:      265.1
## Max.    :     128294      Max.    :     57641.1
## NA's    :      24      NA's    :      24
## county_weights      county_names_all      county_fips_all      imprecise
## Length:33121      Length:33121      Length:33121      Mode :logical
## Class :character      Class :character      Class :character      FALSE:33121
## Mode  :character      Mode  :character      Mode  :character
##
##
##
##      military          timezone
## Mode :logical      Length:33121
## FALSE:33121      Class :character
##                               Mode  :character
##
##
##
##

```

```

for (column in zipcode_data){
  print( typeof(column))
}

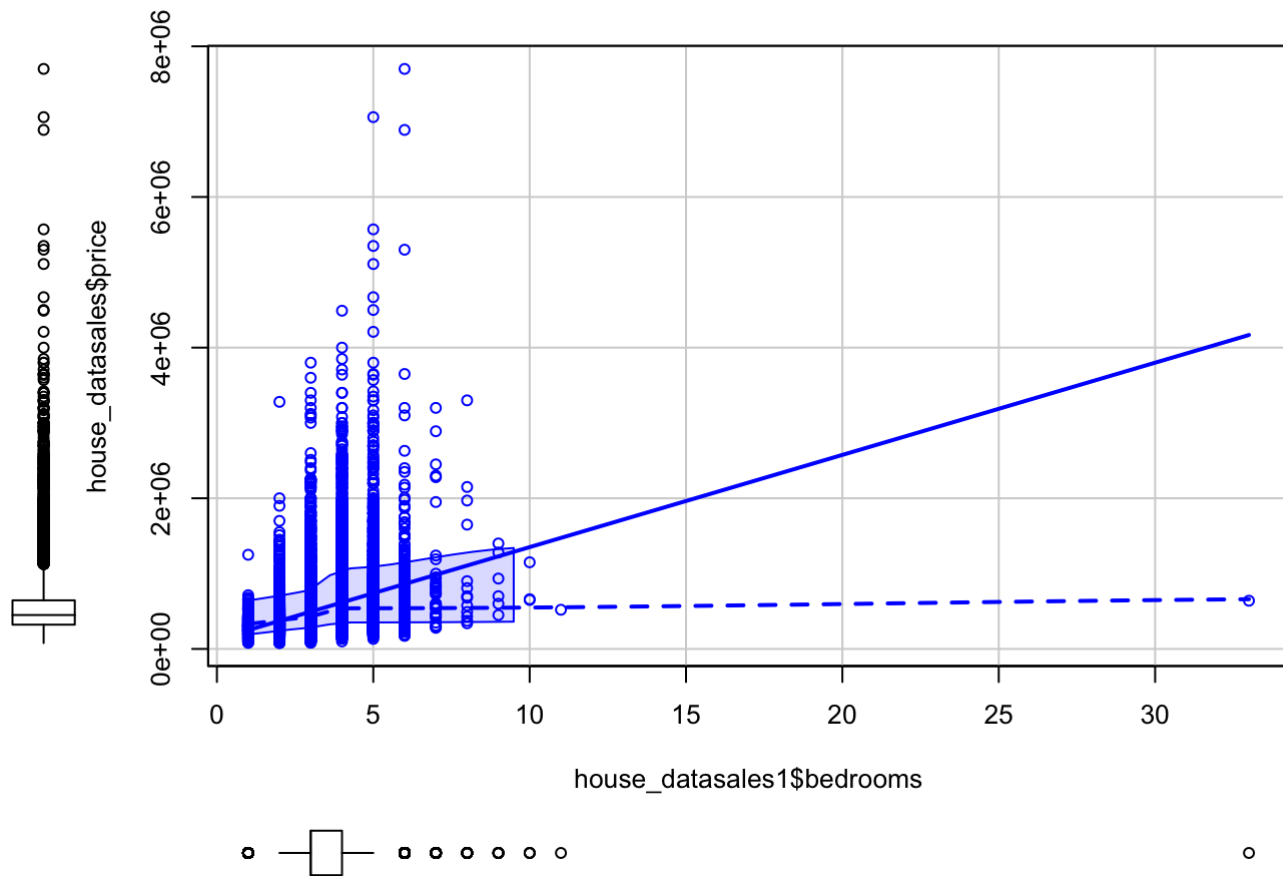
```

```
## [1] "character"
## [1] "double"
## [1] "double"
## [1] "character"
## [1] "character"
## [1] "character"
## [1] "logical"
## [1] "logical"
## [1] "double"
## [1] "double"
## [1] "character"
## [1] "character"
## [1] "character"
## [1] "character"
## [1] "character"
## [1] "logical"
## [1] "logical"
## [1] "character"
```

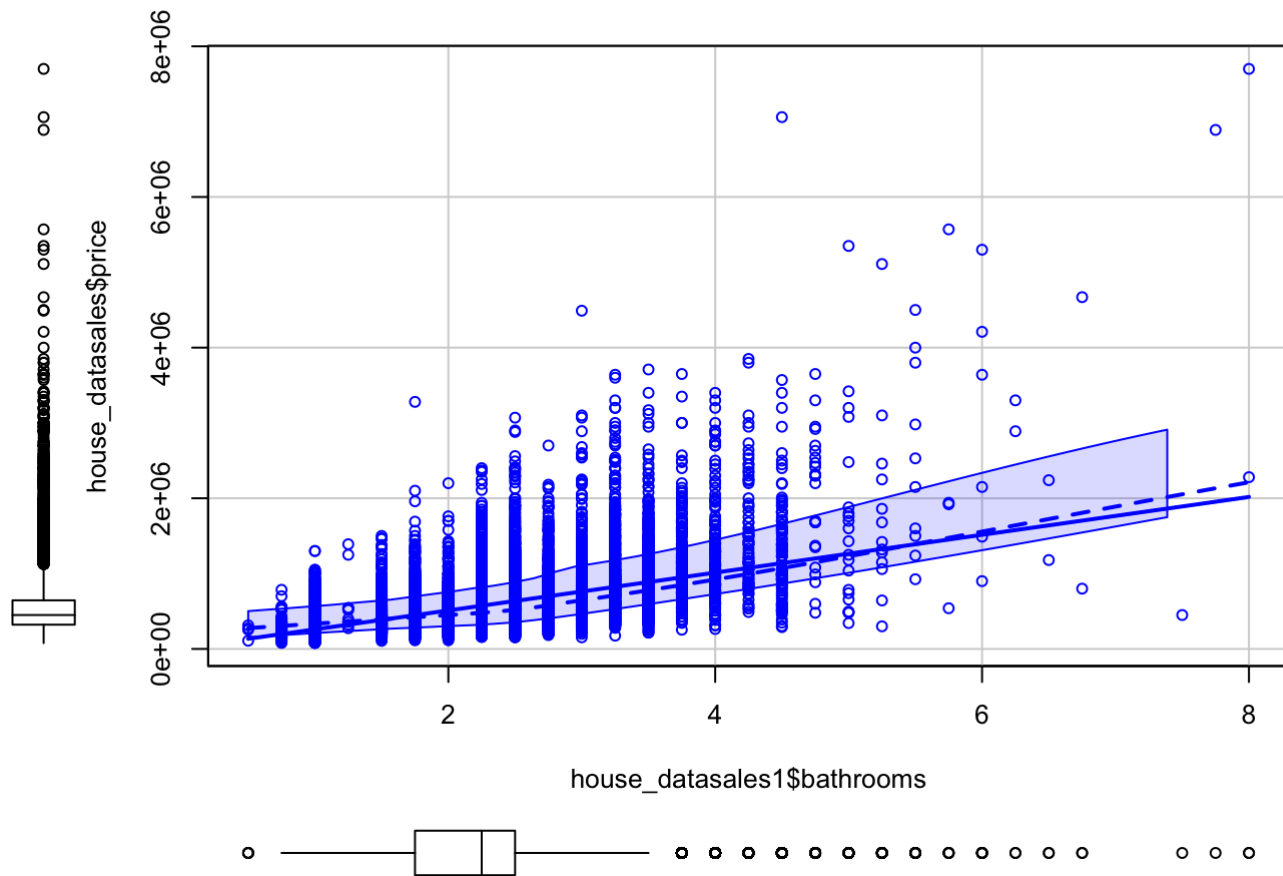
From the above, we understand that 10 features have their data type as character, 4 features have it as double and the remaining 4 are logical.

As our main focus is on prediction of sold price, we remove values that do not have much impact on the change in the value of price.

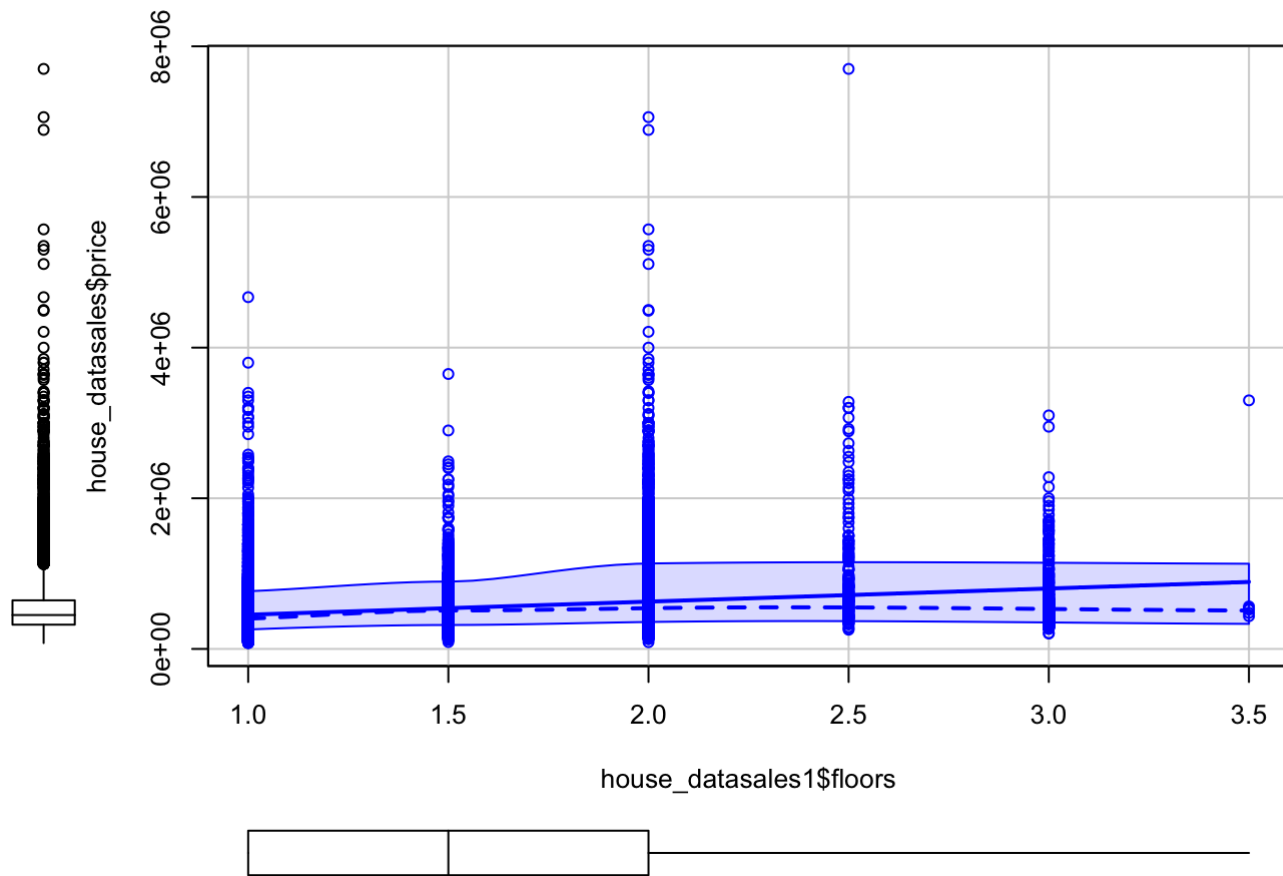
```
par(mfrow=c(4,5))
scatterplot(house_datasales1$bedrooms,house_datasales$price)
```



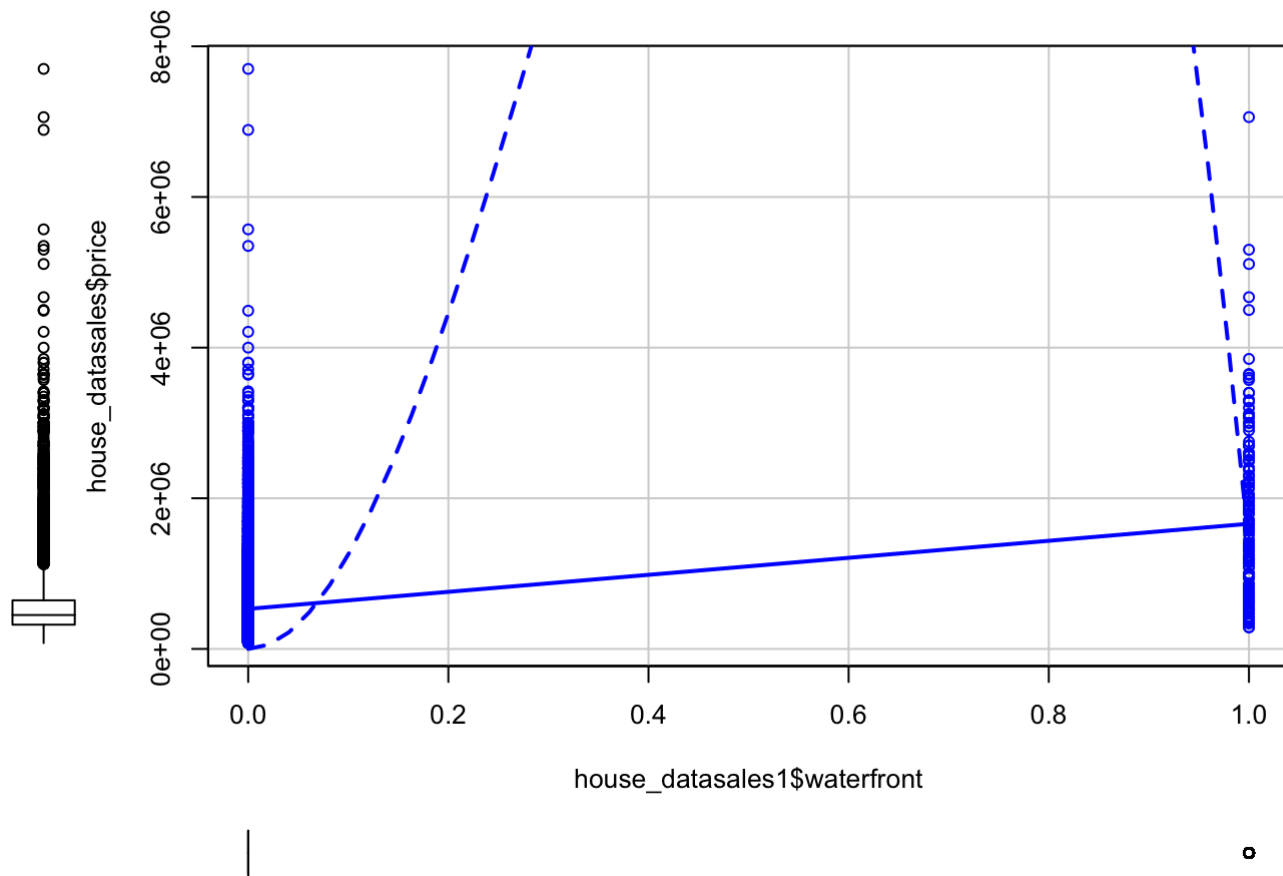
```
scatterplot(house_data$bedrooms,house_data$price)
```



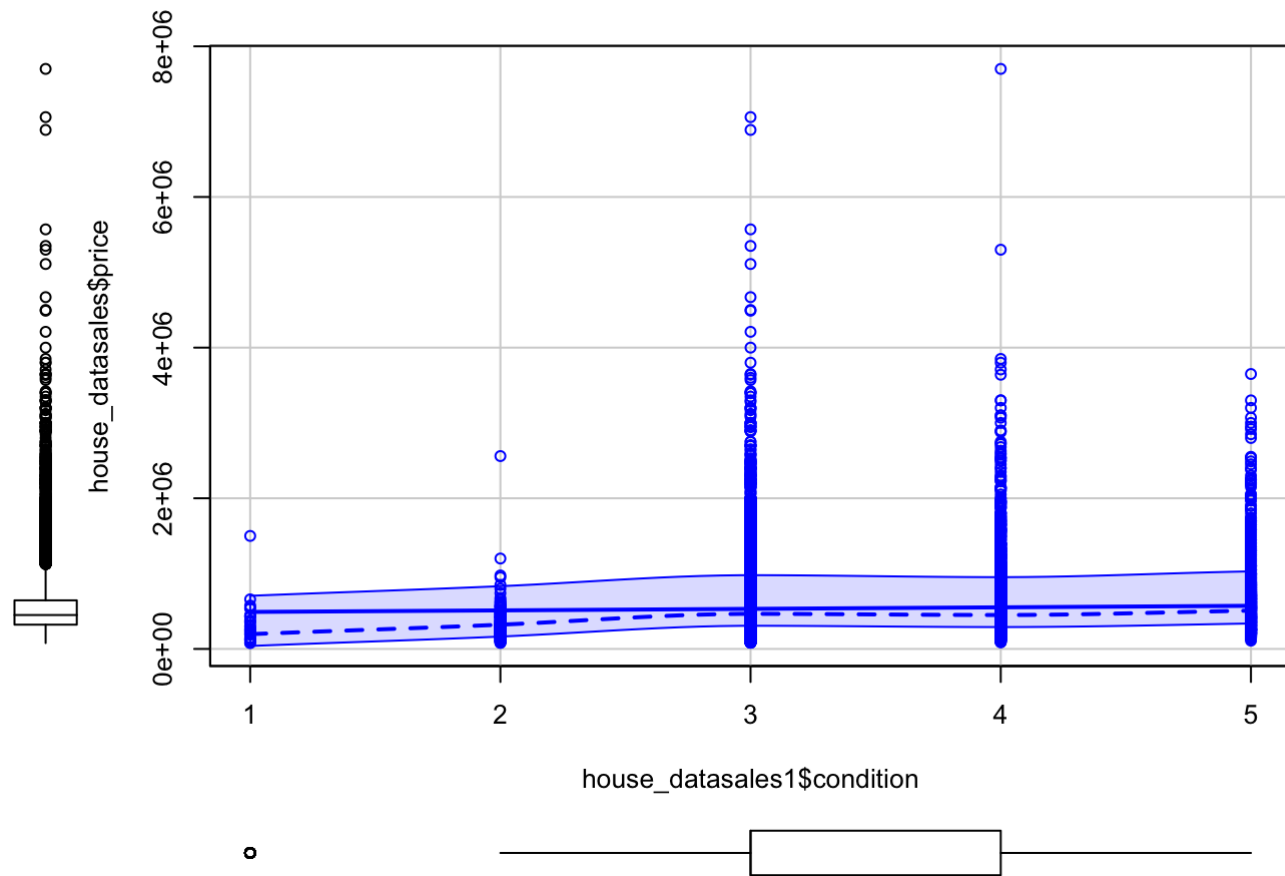
```
scatterplot(house_datasales1$floors,house_datasales$price)
```




```
scatterplot(house_data$waterfront,house_data$price)
```



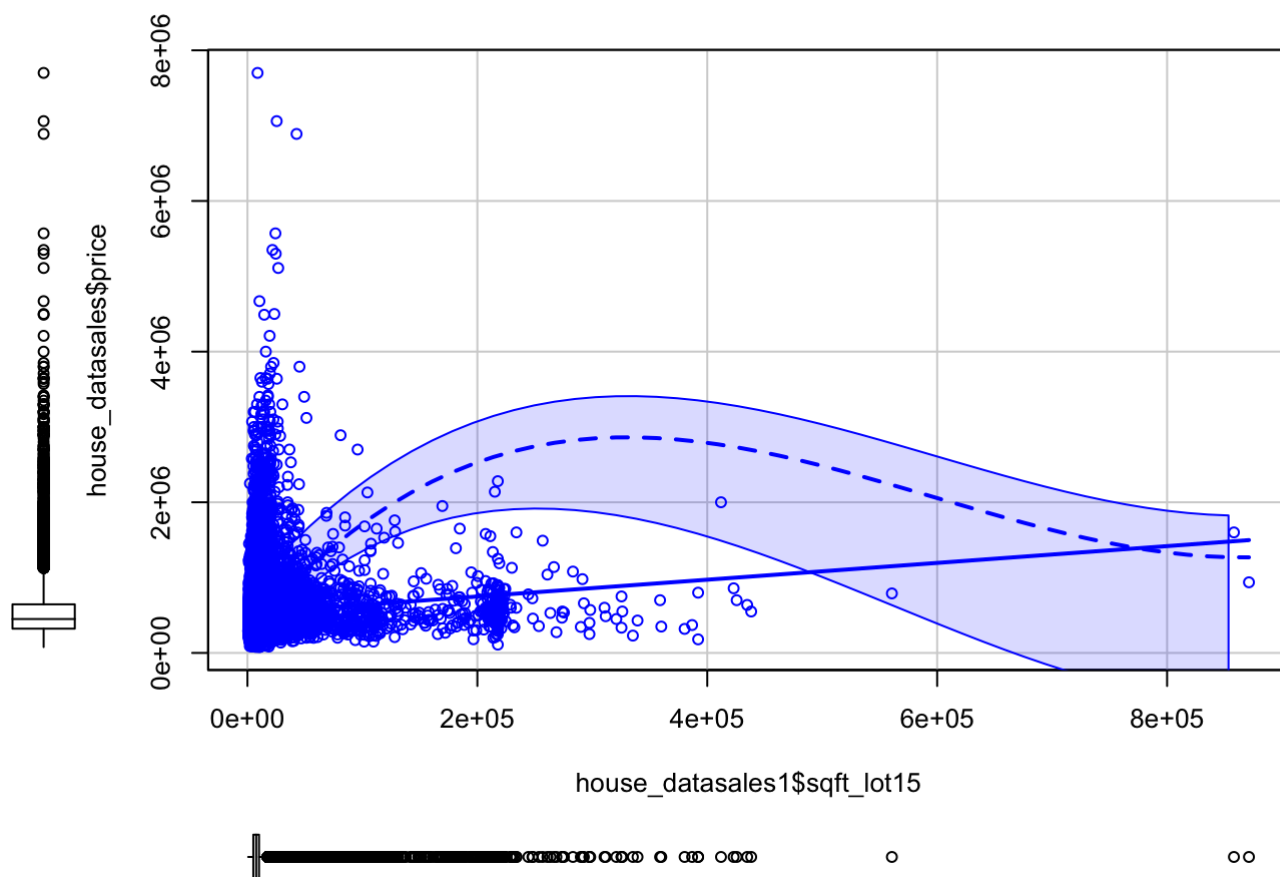
```
scatterplot(house_datsales1$condition,house_datsales$price)
```



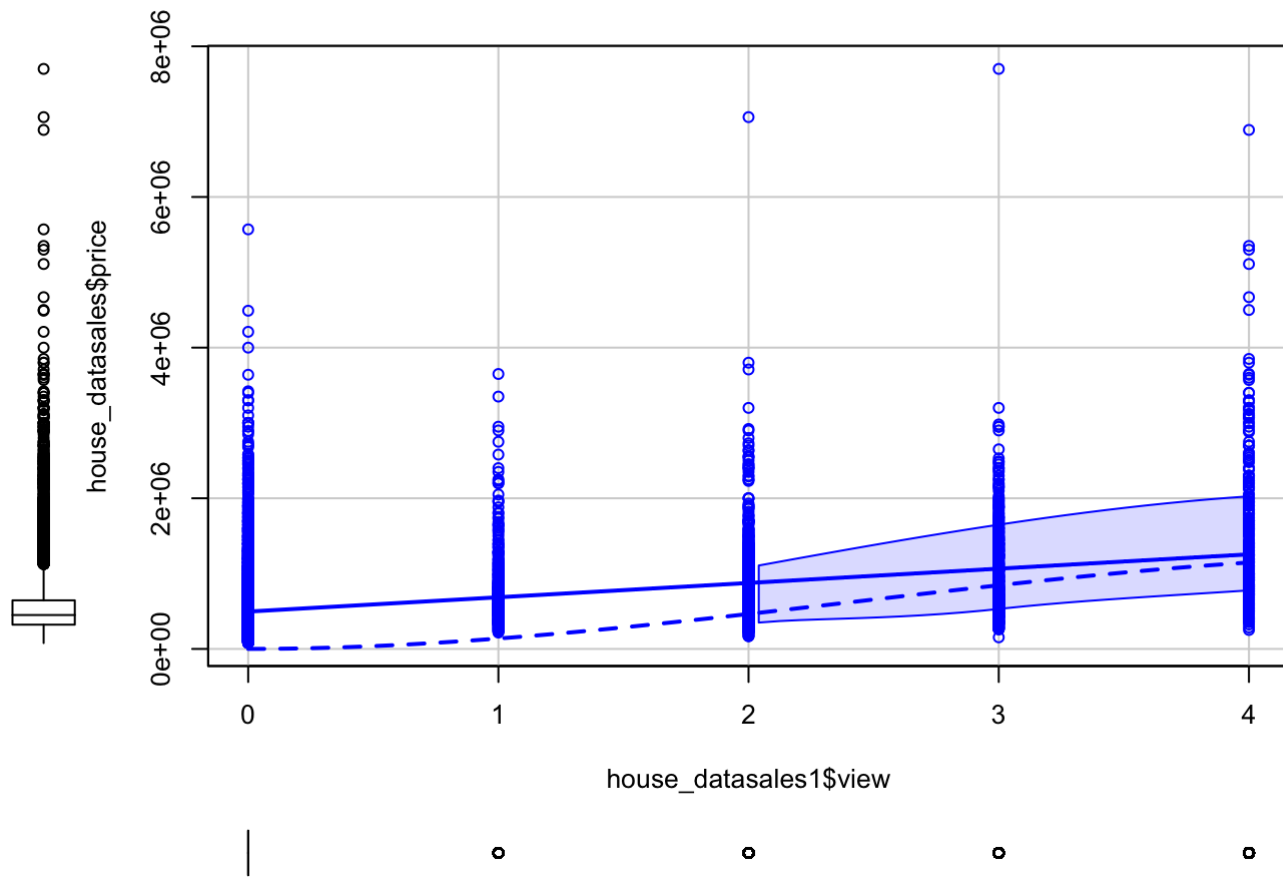
```
scatterplot(house_datasaless$sqft_living15,house_datasaless$price)
```




```
scatterplot(house_datasales1$sqft_lot15,house_datasales$price)
```

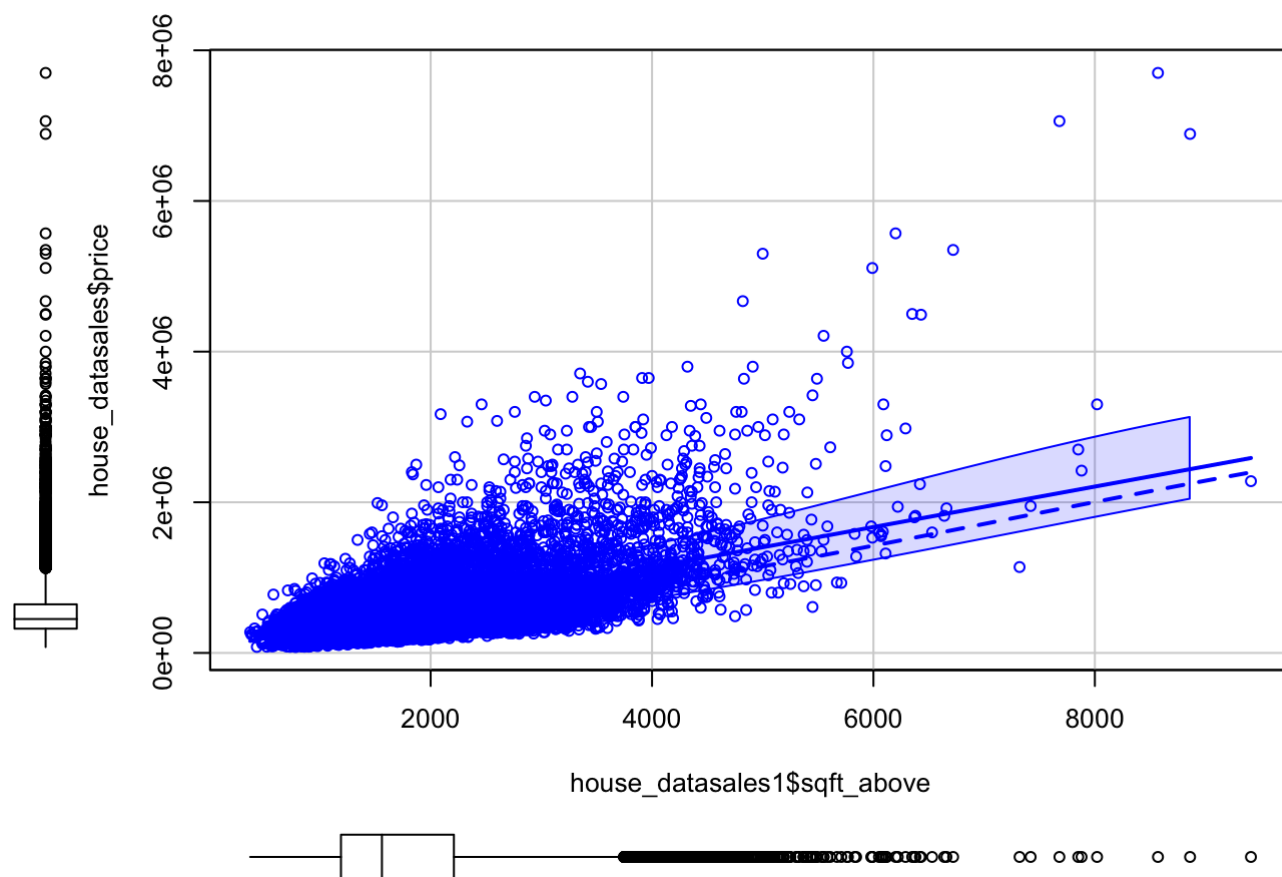
```
scatterplot(house_datasales1$view ,house_datasales$price)
```



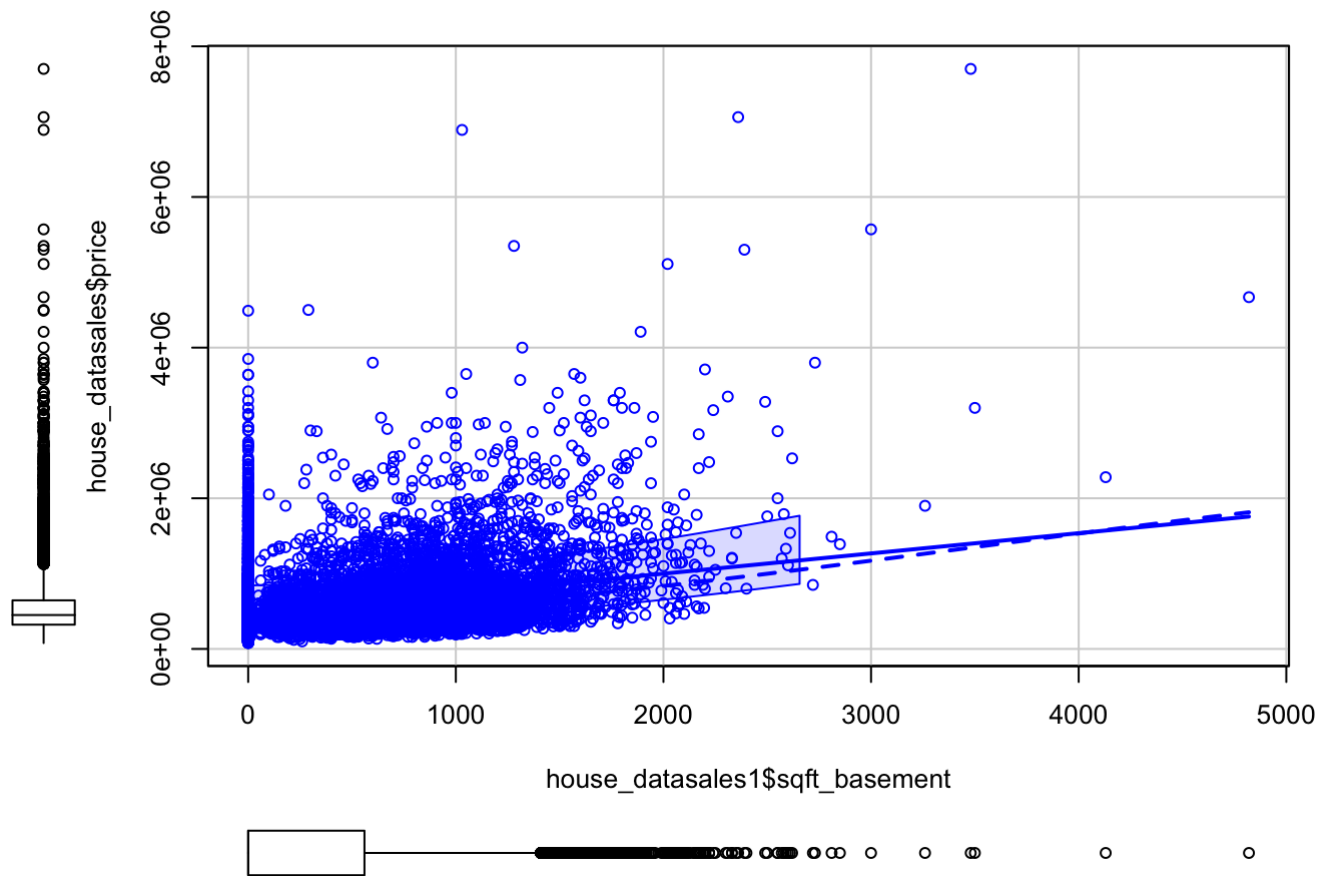
```
scatterplot(house_datasales1$grade,house_datasales$price)
```



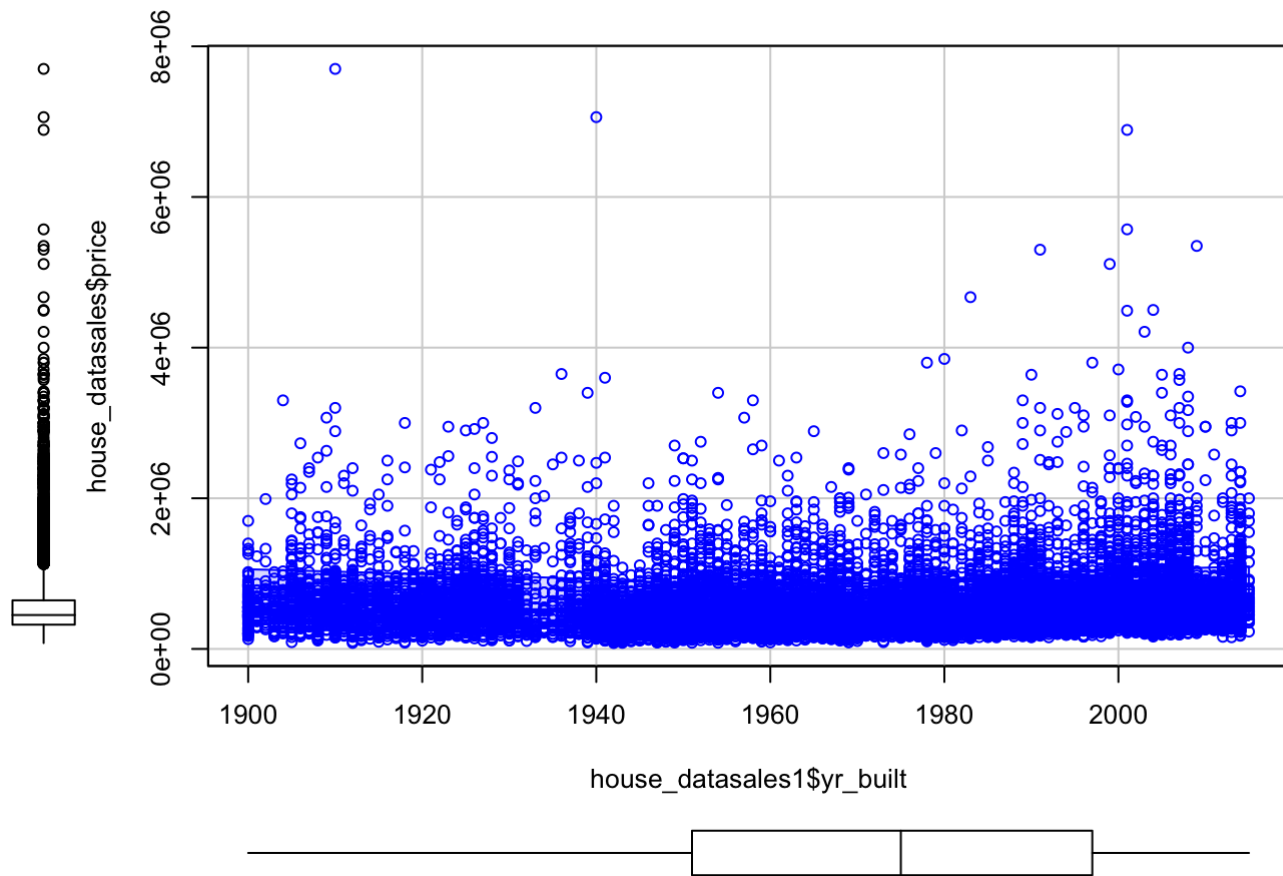
```
scatterplot(house_datasales1$sqft_above,house_datasales$price)
```




```
scatterplot(house_data$ales1$sqft_basement,house_data$ales$price)
```



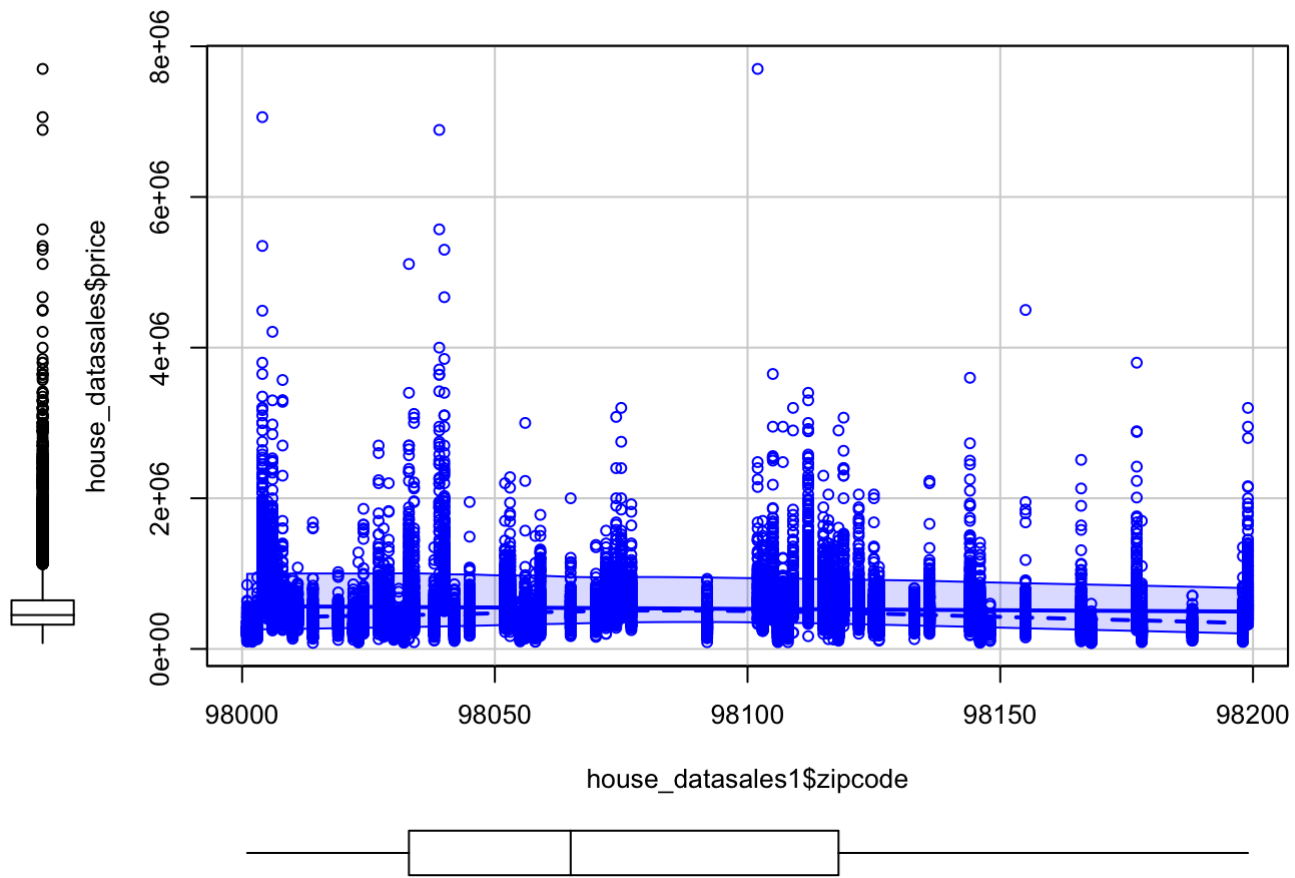
```
scatterplot(house_datasales1$yr_built,house_datasales$price)
```



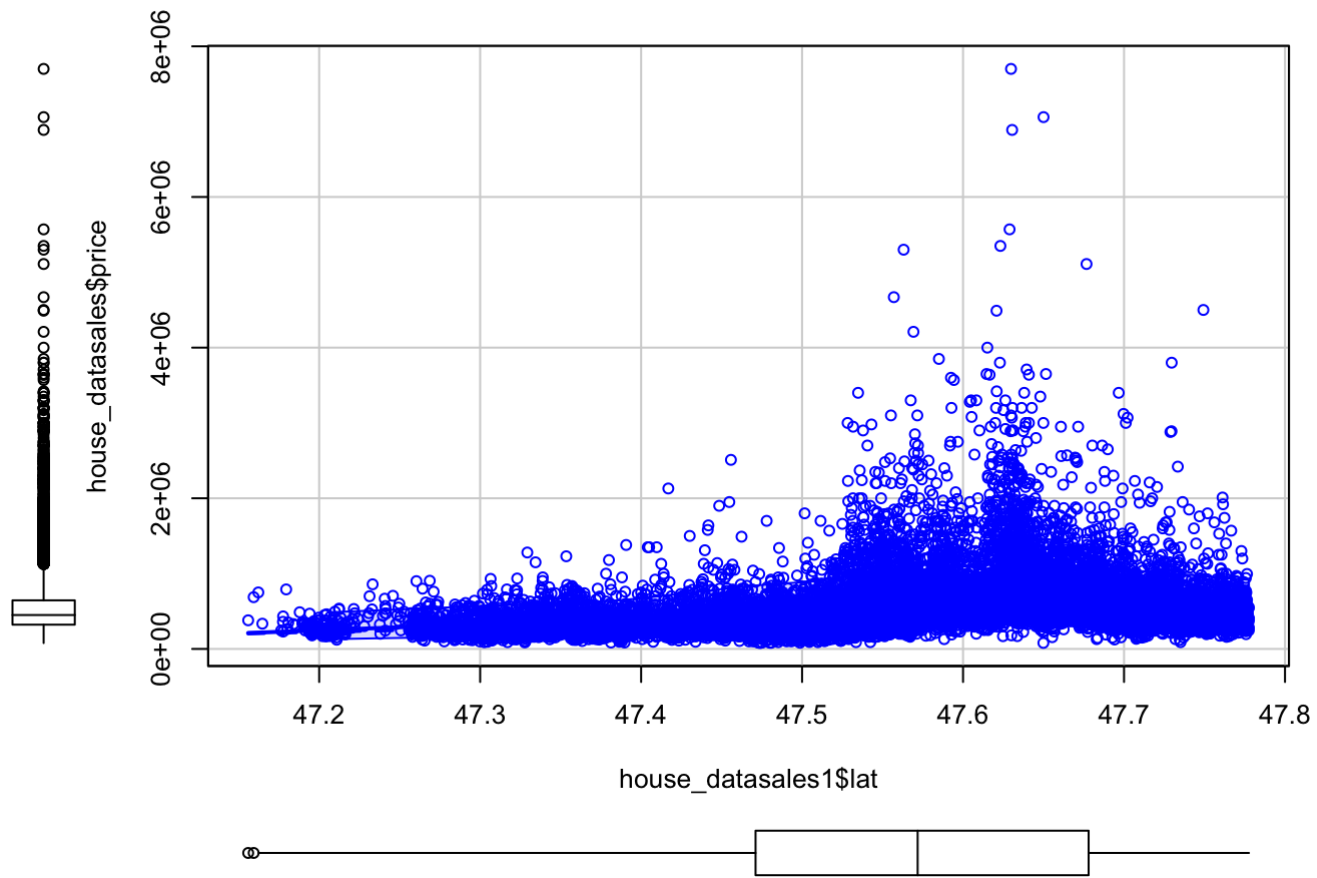
```
scatterplot(house_datasales1$yr_renovated,house_datasales$price)
```




```
scatterplot(house_datasales1$zipcode,house_datasales$price)
```




```
scatterplot(house_sales1$lat,house_sales$price)
```


```
scatterplot(house_datasales1$long,house_datasales$price)
```



```
scatterplot(house_data$long,house_data$price)
```

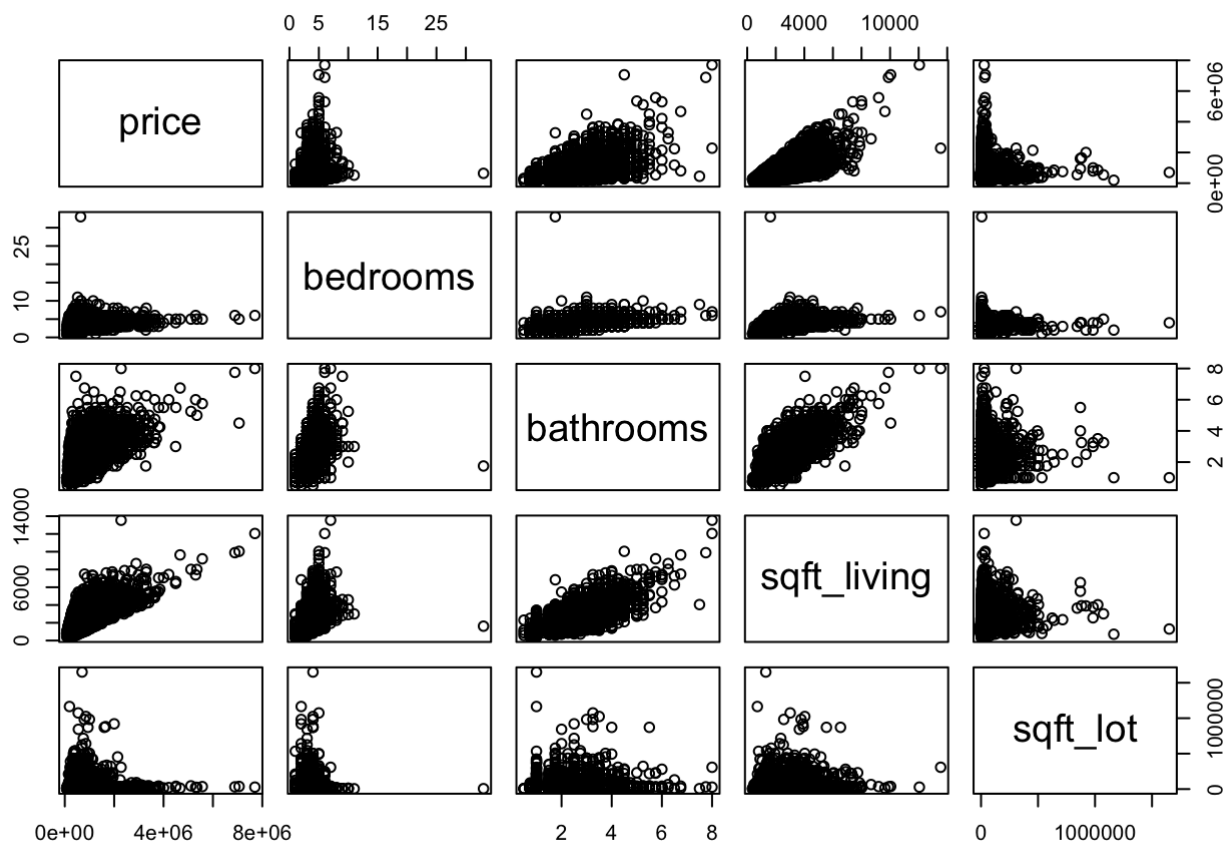


```
scatterplot(house_datasales1$sqft_lot,house_datasales$price)
```

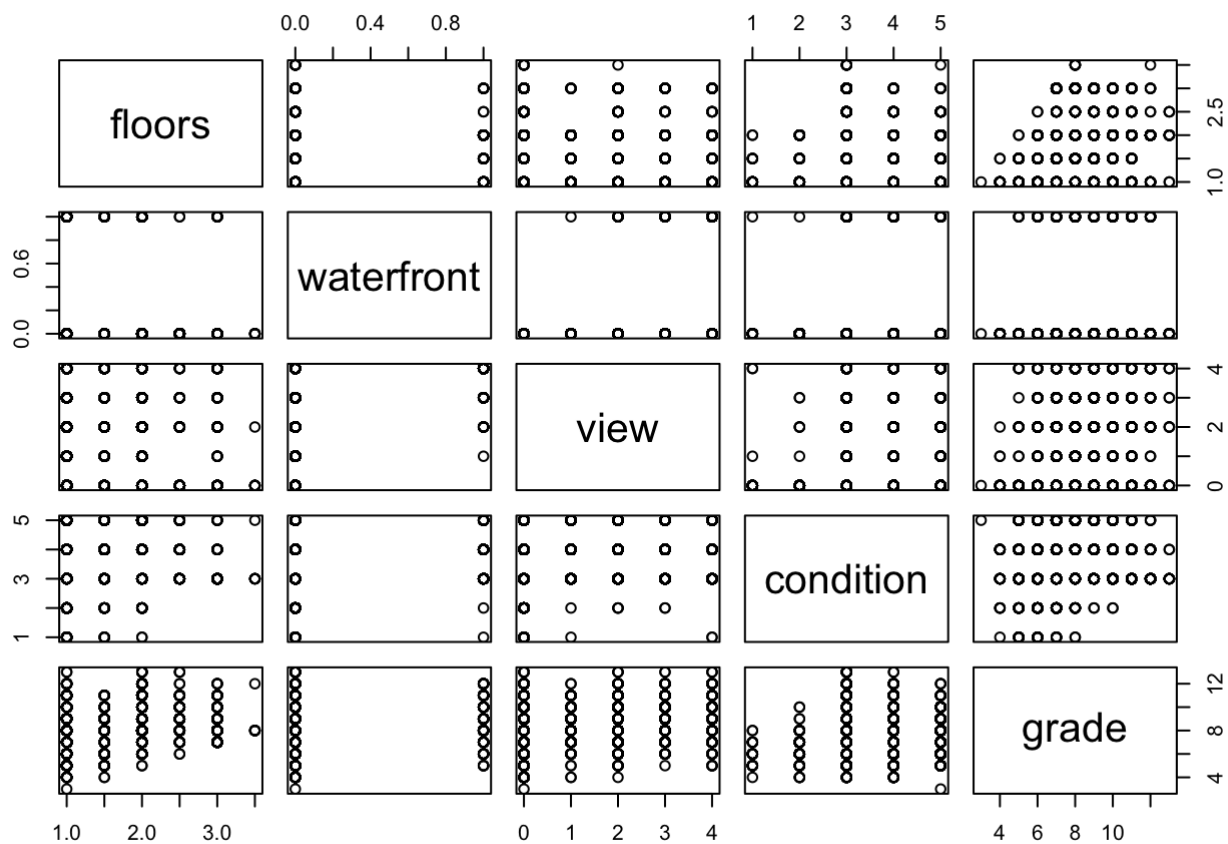


look on the above relation between each variable to the dependent variable - price makes us understand that there are some outliers in the data which we have to take care such that the influence of such points in the creation of the model is less.

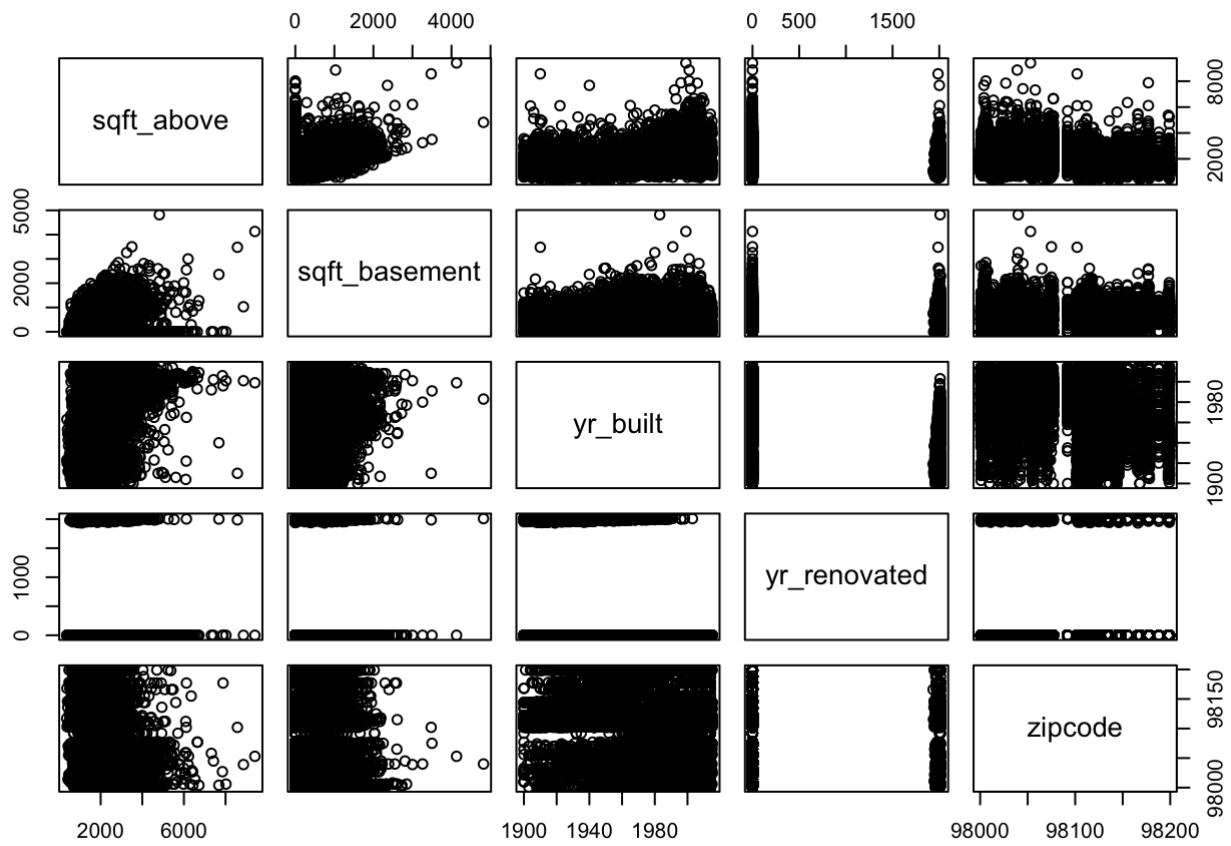
```
plot(house_datasales1[1:5])
```



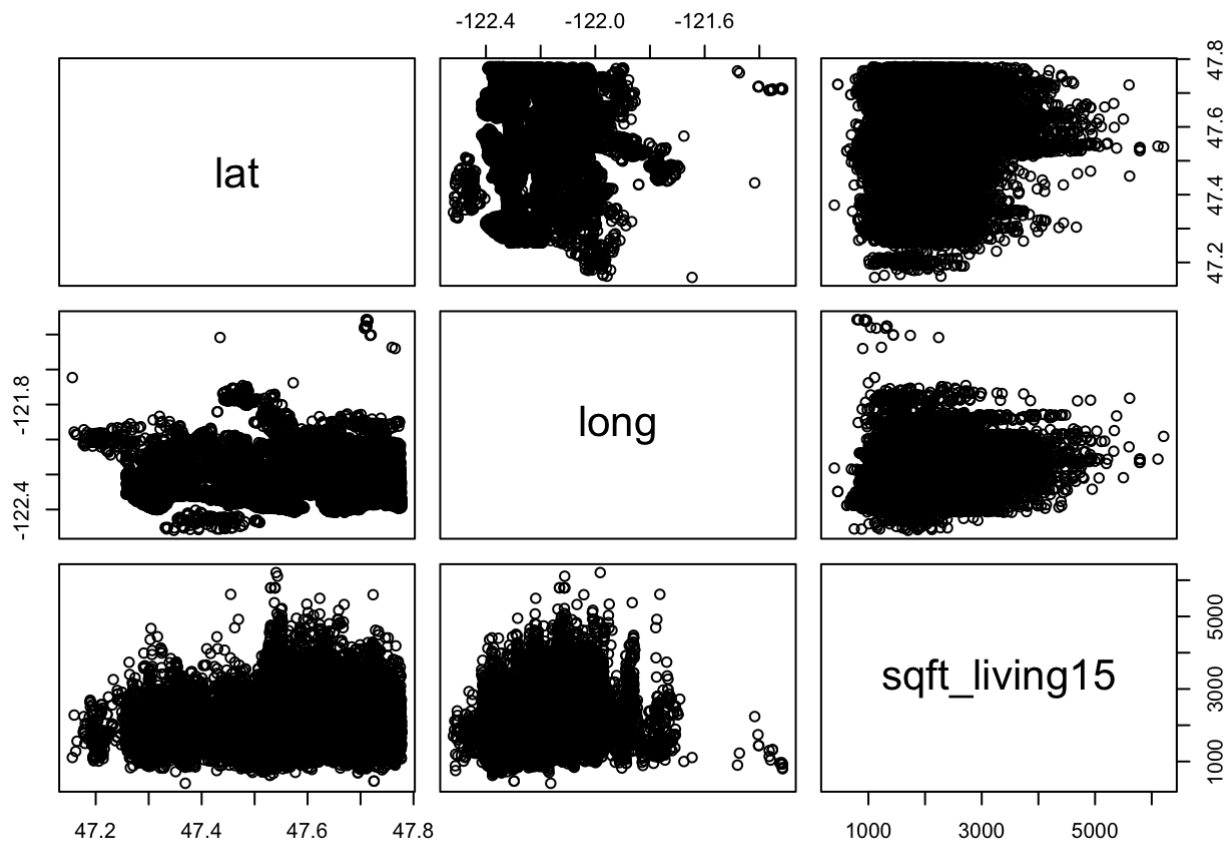
```
plot(house_datasales1[6:10])
```



```
plot(house_datsales1[11:15])
```



```
plot(house_datasales1[16:18])
```



```
cor(house_datsales1[1:5],house_datsales1$price)
```

```
##           [,1]
## price      1.00000000
## bedrooms   0.30878747
## bathrooms   0.52590562
## sqft_living 0.70191730
## sqft_lot    0.08987622
```

```
cor(house_datsales1[6:10],house_datsales1$price)
```

```
##           [,1]
## floors     0.25680354
## waterfront 0.26639846
## view       0.39737030
## condition  0.03605638
## grade      0.66795077
```

```
cor(house_datsales1[11:19],house_datsales1$price)
```



```
##           [,1]
## sqft_above    0.60536794
## sqft_basement 0.32379891
## yr_built      0.05395333
## yr_renovated   0.12642362
## zipcode       -0.05340243
## lat           0.30669231
## long          0.02203632
## sqft_living15  0.58524120
## sqft_lot15     0.08284493
```

```
#View(house_datasales1)
```

Data Preprocessing

Performing Data Sanity Checks before proceeding with analysis

```
house_datasales1$zip <- house_datasales1$zipcode
house_datasales1$zipcode <- NULL
house_datasales1$basement <- house_datasales1$sqft_basement
house_datasales1$sqft_basement <- NULL
#View(zipcode_data)
zipcode_data <- zipcode_data[ -c(2:3,5,7:18) ]
#View(zipcode_data)
## Converting categorical values to numeric
house_datasales1$basement = ifelse(house_datasales1$basement>0,"1","0")
#View(house_datasales1)
house_datasales1$renovation = ifelse(house_datasales1$yr_renovated">0,"1","0")
house_datasales1$yr_renovated <- NULL
```

Checking missing values and duplicate values in the data:

```
## missing values check
print(sum(is.na(house_datasales1)))
```

```
## [1] 0
```

```
print(sum(is.na(zipcode_data)))
```

```
## [1] 0
```

```
## duplicate rows check
zipcode_data %>% distinct(zip, .keep_all= TRUE)
```

```
## # A tibble: 33,121 × 3
##   zip    city      state_name
##   <chr> <chr>      <chr>
## 1 00601 Adjuntas  Puerto Rico
## 2 00602 Aguada    Puerto Rico
## 3 00603 Aguadilla  Puerto Rico
## 4 00606 Maricao    Puerto Rico
## 5 00610 Anasco    Puerto Rico
## 6 00612 Arecibo    Puerto Rico
## 7 00616 Bajadero  Puerto Rico
## 8 00617 Barceloneta Puerto Rico
## 9 00622 Boqueron   Puerto Rico
## 10 00623 Cabo Rojo  Puerto Rico
## # ... with 33,111 more rows
```

```
#View(house_datasales1)
#View(zipcode_data)
final_merged_data <- merge(house_datasales1,zipcode_data,by="zip")
#View(final_merged_data)
```

We go ahead with merging 2 datasets as it will then be easy for us to create the model.

```
# Merging 2 datasets
final_merged_data <- merge(house_datasales1,zipcode_data,by="zip")
#View(final_merged_data)
```

```
# Examine the frequency table of city and state_name
table(final_merged_data$city)
```

```
##
##      Auburn      Bellevue Black Diamond      Bothell      Carnation
##      911        1407        100        195        124
##      Duvall      Enumclaw      Fall City      Federal Way      Issaquah
##      190        233        80        779        733
##      Kenmore      Kent        Kirkland      Maple Valley      Medina
##      283        1201        977        589        50
## Mercer Island      North Bend      Redmond      Renton      Sammamish
##      282        220        977        1597        800
##      Seattle      Snoqualmie      Vashon      Woodinville
##      8973        308        117        471
```


```
table(final_merged_data$state_name)
```

```
##
## Washington
##      21597
```


```
final_merged_data$state_name <- NULL
```

Detecting outliers

```
boxplot(final_merged_data$price, ylab = "Price")
```



```
boxplot(final_merged_data$bedrooms, ylab = "Bedrooms")
```



Exploratory Data Analysis

```
#View(final_merged_data)
summary(final_merged_data)
```

```
##          zip          price          bedrooms          bathrooms
## Min.      :98001   Min.      : 78000   Min.      : 1.000   Min.      :0.500
## 1st Qu.:98033   1st Qu.: 322000   1st Qu.: 3.000   1st Qu.:1.750
## Median :98065   Median : 450000   Median : 3.000   Median :2.250
## Mean      :98078   Mean      : 540297   Mean      : 3.373   Mean      :2.116
## 3rd Qu.:98118   3rd Qu.: 645000   3rd Qu.: 4.000   3rd Qu.:2.500
## Max.      :98199   Max.      :7700000   Max.      :33.000   Max.      :8.000
## sqft_living sqft_lot          floors          waterfront
## Min.      : 370   Min.      : 520   Min.      :1.000   Min.      :0.000000
## 1st Qu.: 1430   1st Qu.: 5040   1st Qu.:1.000   1st Qu.:0.000000
## Median : 1910   Median : 7618   Median :1.500   Median :0.000000
## Mean      : 2080   Mean      : 15099   Mean      :1.494   Mean      :0.007547
## 3rd Qu.: 2550   3rd Qu.: 10685   3rd Qu.:2.000   3rd Qu.:0.000000
## Max.      :13540   Max.      :1651359   Max.      :3.500   Max.      :1.000000
##          view          condition          grade          sqft_above          yr_built
## Min.      :0.0000   Min.      :1.00   Min.      : 3.000   Min.      : 370   Min.      :1900
## 1st Qu.:0.0000   1st Qu.:3.00   1st Qu.: 7.000   1st Qu.:1190   1st Qu.:1951
## Median :0.0000   Median :3.00   Median : 7.000   Median :1560   Median :1975
## Mean      :0.2343   Mean      :3.41   Mean      : 7.658   Mean      :1789   Mean      :1971
## 3rd Qu.:0.0000   3rd Qu.:4.00   3rd Qu.: 8.000   3rd Qu.:2210   3rd Qu.:1997
## Max.      :4.0000   Max.      :5.00   Max.      :13.000   Max.      :9410   Max.      :2015
##          lat          long          sqft_living15          sqft_lot15
## Min.      :47.16   Min.      :-122.5   Min.      : 399   Min.      : 651
## 1st Qu.:47.47   1st Qu.: -122.3   1st Qu.:1490   1st Qu.: 5100
## Median :47.57   Median : -122.2   Median :1840   Median : 7620
## Mean      :47.56   Mean      :-122.2   Mean      :1987   Mean      :12758
## 3rd Qu.:47.68   3rd Qu.: -122.1   3rd Qu.:2360   3rd Qu.:10083
## Max.      :47.78   Max.      :-121.3   Max.      :6210   Max.      :871200
##          basement          renovation          city
## Length:21597   Length:21597   Length:21597
## Class :character   Class :character   Class :character
## Mode :character   Mode :character   Mode :character
##
##
##
```

```
## missing value check
na_check=data.frame(no_of_na_values=colSums(is.na(final_merged_data)))
head(na_check,5)
```

```
##          no_of_na_values
## zip                      0
## price                    0
## bedrooms                 0
## bathrooms                0
## sqft_living              0
```

```
## Sampling the data
set.seed(123)
split = sample.split(final_merged_data$zip, SplitRatio = 0.7)
train =subset(final_merged_data,split == TRUE)
test =subset(final_merged_data, split == FALSE)
dim(train)
```

```
## [1] 15116    20
```


```
#View(train)
dim(test)
```

```
## [1] 6481    20
```

Finding the correlation and plotting the features using heatmap

```
corr_data=data.frame(train[,1:20])
corr_data = corr_data[, -c(18:21)]

correlation=cor(corr_data)
par(mfrow=c(1, 1))
corrplot(correlation,method="color")
```



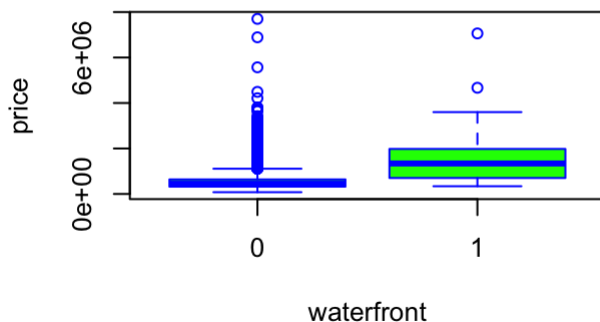
Scatter plots for determining the positive-correlated variables

```
plot1=ggplot(data = train, aes(x = bedrooms, y = price)) +
  geom_jitter() + geom_smooth(method = "lm", se = FALSE)+labs(title="Scatter plot of Price vs Bedrooms", x="Bedrooms",y="Price")
plot2=ggplot(data = train, aes(x = bathrooms, y = price)) +
  geom_jitter() + geom_smooth(method = "lm", se = FALSE)+labs(title="Scatter plot of Price vs Bathrooms", x="Bathrooms",y="Price")
plot3=ggplot(data = train, aes(x = floors, y = price)) +
  geom_jitter() + geom_smooth(method = "lm", se = FALSE)+labs(title="Scatter plot of Price vs Floors", x="Floors",y="Price")
```

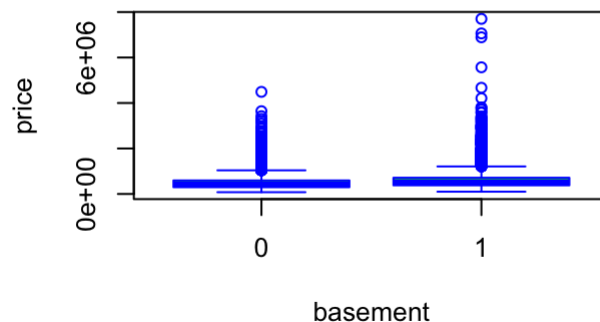
To get clear view of relationships, we plot the boxplots

```
par(mfrow=c(2, 2))
boxplot(price~waterfront,data=train,main="Price vs Waterfront", xlab="waterfront",ylab="price",col="green",border="blue")
boxplot(price ~ basement,data=train,main="Price vs Basement", xlab="basement",ylab="price",col="green",border="blue")
boxplot(price~renovation,data=train,main="Price vs Renovation", xlab="renovation",ylab="price",col="green",border="blue")
boxplot(price~city,data=train,main="Price vs City", xlab="city",ylab="price",col="green",border="blue")
```


Price vs Waterfront



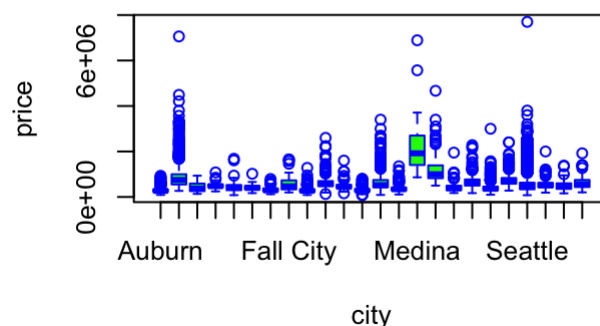
Price vs Basement



Price vs Renovation




Price vs City



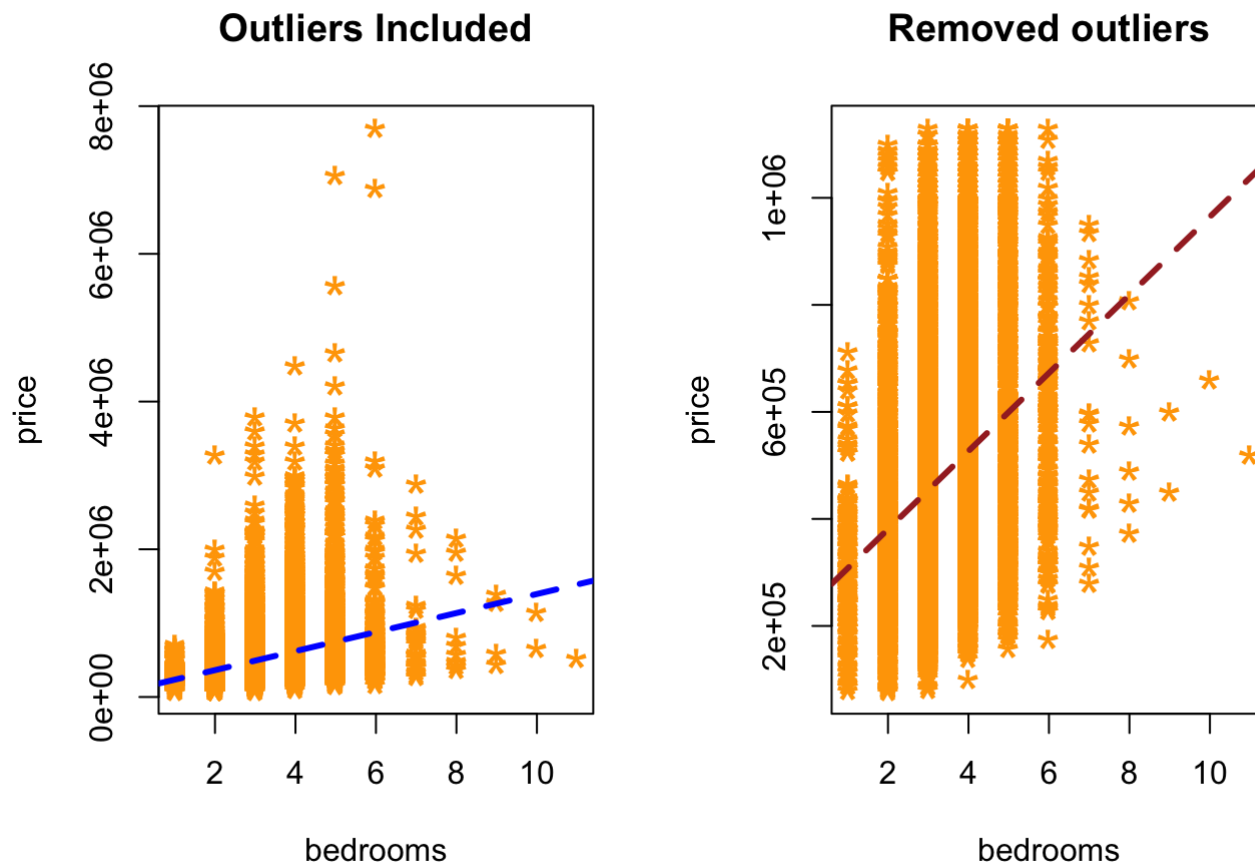
```
ggplot(data=train)+geom_boxplot(aes(x=bedrooms,y=price))
```

```
## Warning: Continuous x aesthetic
## i did you forget `aes(group = ...)`?
```



Plotting data with and without outliers to understand the change in the slope.

```
ol=boxplot(train$price,plot=FALSE)$out
ol_data=train[which(train$price %in% ol),]
train1= train[-which(train$price %in% ol),]
par(mfrow=c(1, 2))
plot(train$bedrooms, train$price, main="Outliers Included", xlab="bedrooms", ylab="price",
     pch="*", col="orange", cex=2)
abline(lm(price ~ bedrooms, data=train), col="blue", lwd=3, lty=2)
plot(train1$bedrooms, train1$price, main="Removed outliers", xlab="bedrooms", ylab="price",
     pch="*", col="orange", cex=2)
abline(lm(price ~bedrooms, data=train1), col="brown", lwd=3, lty=2)
```

Analaysis of variance (ANOVA)


```
## Anova and Turkey test for price vs condition and plotting the distribution
## Calculate frequency, mean and standard deviation
final_merged_data %>% group_by(condition) %>% summarise(condition_freq = n(), price_mean = mean(price, na.rm = TRUE), price_sd = sd(price, na.rm = TRUE))
```

```
## # A tibble: 5 × 4
##   condition condition_freq price_mean price_sd
##   <dbl>         <int>      <dbl>    <dbl>
## 1         1             29   341067.  273483.
## 2         2            170   328179.  246987.
## 3         3          14020   542173.  364650.
## 4         4           5677   521374.  358796.
## 5         5          1701   612578.  411318.
```

```
anova_cond <- aov(price ~ condition, data = final_merged_data)
summary(anova_cond)
```

```
##           Df      Sum Sq   Mean Sq F value   Pr(>F)
## condition     1 3.789e+12 3.789e+12   28.11 1.16e-07 ***
## Residuals 21595 2.911e+15 1.348e+11
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
options(scipen=999)
ggboxplot(final_merged_data, x = "condition", y = "price", ylim=c(78000,7700000))
```




```
## Anova and Turkey test for price vs renovation and plotting the distribution
## Calculate frequency, mean and standard deviation
final_merged_data %>% group_by(renovation) %>% summarise(renovation_freq = n(), price_mean = mean(price, na.rm = TRUE), price_sd = sd(price, na.rm = TRUE))
```

```
## # A tibble: 2 × 4
##   renovation renovation_freq price_mean price_sd
##   <chr>          <int>      <dbl>    <dbl>
## 1 0             20683    530560.  349805.
## 2 1              914    760629.  608017.
```

```
anova_reno <- aov(price ~ renovation, data = final_merged_data)
summary(anova_reno)
```

```
##           Df           Sum Sq          Mean Sq F value           Pr(>F)
## renovation      1  46332107051977 46332107051977    348.8 <0.0000000000000002
## Residuals  21595 2868250023356209   132820098326
##
## renovation ***
## Residuals
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
options(scipen=999)
ggboxplot(final_merged_data, x = "renovation", y = "price", ylim=c(78000,7700000))
```



```
## Anova and Turkey test for price vs city and plotting the distribution
## Calculate frequency, mean and standard deviation
options(dplyr.print_max = 1e9)
final_merged_data %>% group_by(city) %>% summarise(city_freq = n(), price_mean = mean(pr
ice, na.rm = TRUE), price_sd = sd(price, na.rm = TRUE))
```


```
## # A tibble: 24 × 4
##   city          city_freq price_mean price_sd
##   <chr>          <int>     <dbl>    <dbl>
## 1 Auburn           911     291648.  108422.
## 2 Bellevue        1407     898466.  559782.
## 3 Black Diamond   100     423666.  195415.
## 4 Bothell         195     490377.  121971.
## 5 Carnation       124     455617.  258603.
## 6 Duvall          190     424815.  130638.
## 7 Enumclaw        233     316742.  122329.
## 8 Fall City        80     586121.  376719.
## 9 Federal Way     779     289391.  108399.
## 10 Issaquah        733     615122.  260451.
## 11 Kenmore         283     462489.  149530.
## 12 Kent           1201     299470.   91647.
## 13 Kirkland        977     646543.  409633.
## 14 Maple Valley    589     367091.  132721.
## 15 Medina          50     2161300  1166904.
## 16 Mercer Island   282     1194874.  607768.
## 17 North Bend      220     440232.  207554.
## 18 Redmond         977     658432.  231136.
## 19 Renton          1597     403468.  200725.
## 20 Sammamish       800     732821.  280951.
## 21 Seattle        8973     535086.  340519.
## 22 Snoqualmie      308     529630.  185254.
## 23 Vashon          117     489382.  201501.
## 24 Woodinville     471     617498.  244298.
```

```
anova_city <- aov(price ~ city, data = final_merged_data)
summary(anova_city)
```

```
##              Df          Sum Sq      Mean Sq F value          Pr(>F)
## city           23  738104329040975 32091492566999   318.1 <0.0000000000000002
## Residuals    21573  2176477801366957   100888972390
##
## city          ***
## Residuals
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
options(scipen=999)
ggboxplot(final_merged_data, x = "city", y = "price", ylim=c(78000,7700000)) + coord_flip()
```

```
## Coordinate system already present. Adding new coordinate system, which will
## replace the existing one.
```



Data Modelling

Loading the splitted data pre processed Data

Multiple Linear Regression

```
model <- lm(price~bedrooms+bathrooms+floors+waterfront+condition+sqft_living15+sqft_lot1
5+basement+renovation,data=train)
summary(model)
```

```
##
## Call:
## lm(formula = price ~ bedrooms + bathrooms + floors + waterfront +
##      condition + sqft_living15 + sqft_lot15 + basement + renovation,
##      data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1291533  -149571   -25169   103034   5787440
##
## Coefficients:
##              Estimate      Std. Error t value      Pr(>|t|)
## (Intercept)  -455045.81774    16437.09336  -27.684 <0.0000000000000002 ***
## bedrooms      -5625.32457     2905.24193   -1.936    0.0529 .
## bathrooms     101112.98774     4360.25414   23.190 <0.0000000000000002 ***
## floors         53625.05780     5194.68697   10.323 <0.0000000000000002 ***
## waterfront    749134.27427    25185.58577   29.745 <0.0000000000000002 ***
## condition      59809.81332     3480.52171   17.184 <0.0000000000000002 ***
## sqft_living15   235.64690        3.97983   59.210 <0.0000000000000002 ***
## sqft_lot15      -0.27190        0.08391   -3.240    0.0012 **
## basement1      91667.10118     4928.25585   18.600 <0.0000000000000002 ***
## renovation1    201428.42659    10754.29353   18.730 <0.0000000000000002 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 266000 on 15106 degrees of freedom
## Multiple R-squared:  0.4709, Adjusted R-squared:  0.4706
## F-statistic: 1494 on 9 and 15106 DF, p-value: < 0.00000000000000022
```

```
model_fit <- lm(price~bedrooms+bathrooms+floors+waterfront+condition+sqft_living15+sqft_
lot15+basement+renovation, data=train)
s <- stepAIC(model_fit, direction="both")
```

```
## Start:  AIC=377642.2
## price ~ bedrooms + bathrooms + floors + waterfront + condition +
##      sqft_living15 + sqft_lot15 + basement + renovation
##
##              Df          Sum of Sq          RSS      AIC
## <none>                                1068603616353455 377642
## - bedrooms      1      265214779758 10688688831133213 377644
## - sqft_lot15     1      742749127364 1069346365480819 377651
## - floors         1      7538482691711 1076142099045166 377746
## - condition      1     20889274545050 1089492890898504 377933
## - basement       1     24474152529446 1093077768882900 377982
## - renovation     1     24816750138909 1093420366492364 377987
## - bathrooms      1     38041483134714 1106645099488169 378169
## - waterfront     1     62586747317003 1131190363670458 378501
## - sqft_living15  1    248005386066743 1316609002420198 380795
```

```
s$anova
```

```
## Stepwise Model Path
## Analysis of Deviance Table
##
## Initial Model:
## price ~ bedrooms + bathrooms + floors + waterfront + condition +
##      sqft_living15 + sqft_lot15 + basement + renovation
##
## Final Model:
## price ~ bedrooms + bathrooms + floors + waterfront + condition +
##      sqft_living15 + sqft_lot15 + basement + renovation
##
##
##      Step Df Deviance Resid. Df      Resid. Dev      AIC
## 1              15106 1068603616353455 377642.2
```

```
linear_model1 <- lm(price~bedrooms+bathrooms+floors+waterfront+condition+sqft_living15+b
asement+renovation, data=train)
summary(linear_model1)
```

```
##
## Call:
## lm(formula = price ~ bedrooms + bathrooms + floors + waterfront +
##      condition + sqft_living15 + basement + renovation, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1296654 -150346  -25449   102864  5792383
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept)  -456822.205   16433.113  -27.799 <0.0000000000000002 ***
## bedrooms      -5206.009    2903.271   -1.793    0.073 .
## bathrooms     100696.931    4359.733   23.097 <0.0000000000000002 ***
## floors         55046.626    5177.755   10.631 <0.0000000000000002 ***
## waterfront    747541.902    25188.707   29.678 <0.0000000000000002 ***
## condition      59763.067    3481.586   17.165 <0.0000000000000002 ***
## sqft_living15   233.324       3.916   59.583 <0.0000000000000002 ***
## basement1      92843.074    4916.420   18.884 <0.0000000000000002 ***
## renovation1    201291.101   10757.591   18.712 <0.0000000000000002 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 266100 on 15107 degrees of freedom
## Multiple R-squared:  0.4706, Adjusted R-squared:  0.4703
## F-statistic: 1678 on 8 and 15107 DF, p-value: < 0.00000000000000022
```

```
# train the model and store the bootstrap in a dataframe
model_training <- train(price~bedrooms+bathrooms+floors+waterfront+condition+sqft_living
15+basement+renovation, data=train, method="lm")
summary(model_training)
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1296654  -150346   -25449   102864   5792383
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -456822.205    16433.113  -27.799 <0.0000000000000002 ***
## bedrooms      -5206.009     2903.271   -1.793    0.073 .
## bathrooms     100696.931     4359.733   23.097 <0.0000000000000002 ***
## floors         55046.626     5177.755   10.631 <0.0000000000000002 ***
## waterfront    747541.902     25188.707   29.678 <0.0000000000000002 ***
## condition      59763.067     3481.586   17.165 <0.0000000000000002 ***
## sqft_living15    233.324         3.916   59.583 <0.0000000000000002 ***
## basement1      92843.074     4916.420   18.884 <0.0000000000000002 ***
## renovation1    201291.101    10757.591   18.712 <0.0000000000000002 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 266100 on 15107 degrees of freedom
## Multiple R-squared:  0.4706, Adjusted R-squared:  0.4703
## F-statistic: 1678 on 8 and 15107 DF, p-value: < 0.00000000000000022
```

```
model_training_r2 <- summary(model_training$finalModel)$r.squared
model_training_results <- as.data.frame(model_training$results)
```


Influential point Detection using cook's distance

```
cook_distance <- cooks.distance(linear_model1)
sprintf("The mean of Cook's distance is : %f ", mean(cook_distance))
```

```
## [1] "The mean of Cook's distance is : 0.000203 "
```

```
par(mfrow=c(1, 1))
plot(cook_distance, main="i points by Cooks distance")
abline(h = 4*mean(cook_distance, na.rm=T), col="blue")
text(x=1:length(cook_distance)+1,y=cook_distance,labels=ifelse(cook_distance>4*mean(cook
_distance,na.rm=T),names(cook_distance),""), col="blue")
```


i points by Cooks distance



```
i <- as.numeric(names(cook_distance)[(cook_distance > 4*mean(cook_distance, na.rm=T))])
head(train[i, ])
```

##	zip	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	
##	802	98003	225900	3	1.0	1510	8800	1	0
##	1213	98005	1000000	5	2.5	3150	50094	2	0
##	1222	98005	596000	3	2.5	1730	2631	2	0
##	1226	98005	556000	4	2.5	2230	7200	1	0
##	1227	98005	851500	3	2.0	3200	18184	1	0
##	1233	98005	699000	4	2.5	2440	14470	1	0
##	view	condition	grade	sqft_above	yr_built	lat	long	sqft_living15	
##	802	0	4	7	1010	1963	47.3290 -122.330	1290	
##	1213	0	4	9	3150	1969	47.6387 -122.177	3600	
##	1222	0	3	8	1730	2001	47.5878 -122.165	1730	
##	1226	0	4	7	1220	1957	47.5890 -122.156	1920	
##	1227	0	5	8	2000	1977	47.6034 -122.172	1670	
##	1233	0	4	9	1660	1970	47.6401 -122.168	2810	
##	sqft_lot15	basement	renovation	city					
##	802	8470	1	0	Federal Way				
##	1213	48787	0	0	Bellevue				
##	1222	2751	0	0	Bellevue				
##	1226	7200	1	0	Bellevue				
##	1227	7416	1	0	Bellevue				
##	1233	15564	1	0	Bellevue				

```
i_data <- train[i, ]
i_ol <- inner_join(ol_data,i_data)
```

```
## Joining, by = c("zip", "price", "bedrooms", "bathrooms", "sqft_living",
## "sqft_lot", "floors", "waterfront", "view", "condition", "grade", "sqft_above",
## "yr_built", "lat", "long", "sqft_living15", "sqft_lot15", "basement",
## "renovation", "city")
```

```
t2 <- rbind(train,i_ol)
row.names(t2) <- NULL
linear_model2 <- lm(price~bedrooms+bathrooms+floors+waterfront+condition+sqft_living15+b
asement+renovation, data=t2)
summary(linear_model2)
```

```
##
## Call:
## lm(formula = price ~ bedrooms + bathrooms + floors + waterfront +
##      condition + sqft_living15 + basement + renovation, data = t2)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-1298667	-150507	-25503	103345	5787812





```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-457734.108	16440.466	-27.842	<0.0000000000000002 ***
bedrooms	-5671.758	2903.978	-1.953	0.0508 .
bathrooms	101452.341	4358.815	23.275	<0.0000000000000002 ***
floors	54386.975	5178.322	10.503	<0.0000000000000002 ***
waterfront	748345.795	25097.016	29.818	<0.0000000000000002 ***
condition	59751.564	3483.271	17.154	<0.0000000000000002 ***
sqft_living15	234.441	3.909	59.970	<0.0000000000000002 ***
basement1	92803.861	4917.781	18.871	<0.0000000000000002 ***
renovation1	200856.736	10765.092	18.658	<0.0000000000000002 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 266300 on 15119 degrees of freedom
## Multiple R-squared:  0.473, Adjusted R-squared:  0.4727
## F-statistic: 1696 on 8 and 15119 DF, p-value: < 0.00000000000000022
```

Model Evaluation

```
## regression diagnostics
par(mfrow = c(2, 2))
plot(linear_model2)
```



```
## multicollinearity test
## this shows there is no multicollinearity in the model.
vif(linear_model2)
```

	bedrooms	bathrooms	floors	waterfront	condition
##	1.465613	2.424348	1.656469	1.022634	1.096146
	sqft_living15	basement	renovation		
##	1.550625	1.235218	1.022778		

```
## accuracy
prediction_test=predict(newdata=test, linear_model2)
actual_model_fitted_test=data.frame(actual=test$price, predicted=prediction_test)
abs_diff_test = mean(abs(actual_model_fitted_test$actual-actual_model_fitted_test$predicted)/actual_model_fitted_test$actual)
accuracy=1-abs_diff_test
sprintf(" The accuracy of the prediction on test data is : %f",accuracy*100)
```

```
## [1] " The accuracy of the prediction on test data is : 63.784002"
```